

# Soft-Impute Matrix Completion Project Report

Dapeng Hu, Qinglong Tian, Haozhe Zhang, Min Zhang  
Department of Statistics, Iowa State University, Ames, IA 50011

## 1 Summary

The purpose of matrix completion is to solve the problems of the following type: given data in the form of an  $m \times n$  matrix  $Z = \{Z_{ij}\}$ , find an approximating matrix  $\hat{Z}$  that imputes or fills in missing entries in  $Z$ . There are wide applications of matrix completion. One popular example is the movie-rating problem in the "Netflix" Competition, where the data is the basis for a recommender system.

The matrix completion problem is ill-specified unless we impose additional constraints on the unknown matrix  $Z$ , and one common choice is a rank constraint. Suppose that we observe all entries of the matrix  $Z$  indexed by the subset  $\Omega \subset \{1, \dots, m\} \times \{1, \dots, n\}$ . One approach is to consider estimators based on optimization problems of the form

$$\hat{Z} = \arg \min_M \|Z - M\|_F^2 \quad \text{subject to } \text{rank}(M) \leq \delta, \quad (1)$$

where  $\|\cdot\|_F^2$  is Frobenius norm of a matrix. However, this rank-minimization problem is computationally intractable (NP-hard), and cannot be solved in general even for moderately large matrices. Alternatively, we consider the following small modification to (1):

$$\hat{Z} = \arg \min_M \|Z - M\|_F^2 \quad \text{subject to } \|M\|_* \leq \delta, \quad (2)$$

where  $\|\cdot\|_*$  is the nuclear norm, i.e., the sum of the singular values. Under many situations the nuclear norm is an effective convex relaxation to the rank constraint, and hence problem (2) is convex.

Soft-Impute Algorithm was proposed and studied in [5] to solve problem (2). It was shown in [5] that Soft-Impute Algorithm is guaranteed to converge at least sub-linearly, meaning that  $O(\frac{1}{\epsilon})$  iterations are sufficient to compute a solution that is  $\epsilon$ -close to the global optimum. The procedure for Soft-Impute Algorithm is introduced briefly as follows. First consider the case where there is no missing data, to solve (2), we simply compute the SVD of  $Z$ , soft-threshold the singular values by  $\lambda$ , and reconstruct the matrix. The singular value decomposition (SVD) provides an effective method here. Then, we start with an initial guess for the missing values, compute the (full rank) SVD, and then soft-threshold its singular values by  $\lambda$ . We reconstruct the corresponding SVD approximation and obtain new estimates for the missing values. This process is repeated until convergence.

## 2 Various SVD Implementations and Comparison

### 2.1 Brief introduction to SVD Implementations

- **C code via LAPACK:**

The core part of the c code for implementing SVD is the LAPACK routine DGESVD and DGEMM. The first argument of the C main function is filename, and the second argument is  $\lambda$  used in Soft-Impute Algorithm. The format of the input data is like the given Lena dataset. The first row is the dimension of the matrix with N is the row, and P is the column.  $N \neq P$ . The second row is the matrix arranged by column.

- **R internal svd function:**

The main functions used are the LAPACK routines DGESDD and ZGESDD. The performance of R internal svd function essentially depends on the performance of LAPACK.

- **svd Package:**

There are three functions in the svd package that are related to the singular-value decomposition. Because `trlan.svd` and `ztrlan.svd` will not return the right singular vectors and we need that vector in soft-impute algorithm. So we are not going to use those two functions. We only use `propack.svd` to implement the soft-impute algorithm. PROPACK does SVD via the implicitly restarted Lanczos bidiagonalization with partial reorthogonalization.

- **RcppArmadillo Package:**

The RcppArmadillo package includes the header files from the Armadillo library, which is a templated C++ linear algebra library. Various matrix decompositions are provided through optional integration with LAPACK and ATLAS libraries<sup>1</sup>. We use svd in Armadillo to implement the soft-impute algorithm.

- **irlba package:**

We use `irlba` function in the `irlba` package. The augmented implicitly restarted Lanczos bidiagonalization algorithm (IRLBA) finds a few approximate largest singular values and corresponding singular vectors of a sparse or dense matrix using a method of Baglama and Reichel. It is a fast and memory-efficient way to compute a partial SVD. The maximum number of singular values that can be calculated in this function is where is the dimension of the matrix. To find the suitable number of singular values in each iteration of the soft-impute algorithm, we compare the number of thresholded singular values with the maximum allowed number of singular values then take the minimum of these two values as the suitable number in each iteration.

### 2.2 Simulation Study

There are many aspects that can reflect performance of a method such as training error, speed and memory usage. In our comparison, these four methods are just different implementations of the same algorithm so they should have the same training error. And in our simulation

experiments, they do have the same training error so we are not going to compare the training error. We are going to compare the speed of each method in our experiments.

By learning the soft-impute algorithm, we think that three factors may influence the speed. They are dimension of the matrix, missing rate and true rank of the matrix. In our experiment, we let one variable vary and keep other two variables the same to say how each variable affect the performance.

To compare the speed of each method, we use microbenchmark to calculate the time usage and take the median as the result.

### 2.2.1 Comparison under different matrix dimensions

Table 1: Time consumption under different matrix dimensions with missing rate = 0.5 and true matrix rank = 8

No. of rows	No. of columns	R interval svd	propack.svd	RcppArmadillo	irlba
20	20	0.00246	0.00513	<b>0.00232</b>	0.00335
50	50	0.02001	0.03667	0.01876	<b>0.01680</b>
100	100	0.12742	0.20482	<b>0.12269</b>	0.13500
500	500	18.0418	28.8572	<b>17.0435</b>	72.1405
1000	1000	<b>154.7145</b>	237.7374	164.2245	1197.6262

In general, the result shows that the computation speed of R internal svd and RcppArmadillo methods are similar and are the fastest method. When the matrix size is relatively large (1000,1000), the R internal svd performs a little better than RcppArmadillo. Otherwise, RcppArmadillo is a little better.

The propack.svd method is slower than the two methods mentioned above. While the relative difference with those two becomes less and less as the dimension becomes larger and larger.

As for the irlba method, when the dimension is relatively small ( $\leq 100$ ), this method performs similarly compared with R internal svd and RcppArmadillo methods. But when the matrix dimension is relatively large, the time usage of this method will go up dramatically as we can see in Table 1.

### 2.2.2 Comparison under different missing rates

As a whole, the R internal svd and RcppArmadillo perform similarly while RcppArmadillo is a little bit faster than R internal svd.

The propack.svd method is the most unstable method in this situation. When the missing rate is too large or too small, this method fails to execute. Even though the missing rate is suitable, this method is also the slowest one among all the methods.

irlba method performs better as the missing rate goes up. For other three methods, they use more time as the missing rate goes up. While the irlba method is different.

When the missing rate is large ( $\geq 0.6$ ) and the dimension is medium (e.g.,  $100 \times 100$ ), the irlba method is the fastest method.

Table 2: Time consumption under different missing rates with matrix dimension =  $100 \times 100$  and true matrix rank = 8

Missing rate	R interval svd	propack.svd	RcppArmadillo	irlba
0.4	0.10531	*	<b>0.10293</b>	0.13194
0.5	0.13925	0.21978	<b>0.13327</b>	0.14786
0.6	0.14494	0.23622	0.14057	<b>0.12816</b>
0.7	0.19295	0.30568	0.18452	<b>0.13670</b>
0.9	0.28146	*	0.26986	<b>0.09553</b>

### 2.2.3 Comparison under different matrix ranks

Table 3: Time consumption under different matrix ranks with matrix dimension =  $100 \times 100$  and missing rate = 0.5

True matrix rank	R interval svd	propack.svd	RcppArmadillo	irlba
4	0.11331	0.18107	0.10620	0.10551
8	0.12308	0.19143	0.11715	0.12893
12	0.14849	0.23911	0.14536	0.16806
16	0.16822	*	0.17642	0.23704

All of the methods use more time as the true rank goes up. RcppArmadillo method is the fastest although it is just a little faster than R internal svd method. The propack.svd is still the slowest and most unstable method because when the true rank is 16 it fails to execute again. For the irlba method, when the true rank is small ( $eg \leq 8$ ), the speed is very competitive. When the true rank is large, it becomes slower than R internal svd and RcppArmadillo method.

## 3 Application

### 3.1 Lena Image Data

After comparing different methods, we decide to use R internal svd to solve this problem. To design a grid of , we first use svd to decompose the training matrix and get the singular values of it . Then we use the quantile of the singular values of the training matrix to design the grid of . There are two reason why we do in this way. The first is that we have no idea about the range of so we need some reference to specify. The second reason is that by using quantile we can reduce the number of search times. We first use (The subscript represents the quantile) as the vector. And calculate the RMSE in the validation set and we get the Figure 1.

We find that the minimal of the RMSE corresponding to . Then we do further search. We use as the second vector. After calculating the RMSE in the validation set we get the Figure 2. This time we find the minimal of the RMSE corresponding to . We continue our search to find the best . We use as the third vector. The RMSE can be seen in the Figure




Figure 1




Figure 2




Figure 3

3. In Figure 3, we can see that the minimal of RMSE corresponding to which is the best we choose. As we can see, by using quantile, we only need  $s$  to select the turning parameter. In this way, we can reduce the total number of  $s$  to be considered and therefore save a lot of time. Figure 4 is the image with 40% pixels randomly missing. And Figure 5 is the image after using the soft-impute to fix. The total rank in the reconstructed image is 83.

### 3.2 MovieLens Dataset

## 4 Accelerating Soft-Impute

Let us consider a general setting of convex composite minimization problem

$$\min_{\mathbf{x}} F(\mathbf{x}); = f(\mathbf{x}) + g(\mathbf{x}),$$

where  $f(\mathbf{x})$  is convex and  $L$ -smooth (i.e., gradient is Lipschitz continuous), and  $g(\mathbf{x})$  is convex and non-smooth. The proximal operation of  $g$  is easy to calculate here. The proximal gradient method [2] is used to solve this problem: for each iteration  $k = 0, 1, 2, \dots$ ,

$$\mathbf{x}_{k+1} = \text{prox}_{\gamma_k, g}(\mathbf{x}_k - \gamma_k \nabla f(\mathbf{x}_k)) = \arg \min_{\mathbf{x}} \left[ \frac{1}{2\gamma_k} \|\mathbf{x} - \{\mathbf{x}_k - \gamma_k \nabla f(\mathbf{x}_k)\}\|_2^2 + g(\mathbf{x}) \right]$$

The following theorem shows that the proximal gradient method attains the convergence rate of  $O\left(\frac{1}{k}\right)$ .

**THEOREM 1** *Proximal gradient method with fixed step size  $\gamma_k = \frac{1}{L}$  satisfies the following convergence rate*

$$F(\mathbf{x}_k) - F(\mathbf{x}^*) \leq \frac{L \|\mathbf{x}_0 - \mathbf{x}^*\|_2^2}{2k}.$$

Originally developed in [6] and [1], we can accelerate the proximal gradient method as follows:

$$\begin{aligned} \mathbf{x}_{k+1} &= \text{prox}_{\gamma_k, g}(\mathbf{y}_k - \gamma_k \nabla f(\mathbf{y}_k)) \\ \mathbf{y}_{k+1} &= \mathbf{x}_{k+1} + \beta_k (\mathbf{x}_{k+1} - \mathbf{x}_k). \end{aligned}$$



**Figure 4:** Imputed Lena Image



**Figure 5:** Original Lena Image

A common choice for  $\beta$ :  $\beta_k = \frac{\lambda_k - 1}{\lambda_k + 1}$  where  $\lambda_0 = 0$ ,  $\lambda_{k+1} = \frac{1 + \sqrt{1 + 4\lambda_k^2}}{2}$ . The acceleration for this choice of  $\beta$  is called the Fast Iterative Soft Thresholding Algorithm (FISTA) derivation [1]. It was shown in [1] that

**THEOREM 2** *The sequences  $\mathbf{x}_k$ ,  $F(\mathbf{x}_k)$  generated via FISTA with either a constant or back-tracking (with ratio  $\alpha \geq 1$ ) stepsize rule satisfy*

$$F(\mathbf{x}_k) - F(\mathbf{x}^*) \leq \frac{2\alpha L \|\mathbf{x}_0 - \mathbf{x}^*\|_2^2}{k^2}.$$

Soft-Impute Algorithm is a special example of proximal gradient method:

$$\min_Z \underbrace{\frac{1}{2} \|W - Z\|_F^2}_{f(Z)} + \underbrace{\lambda \|Z\|_*}_{g(Z)}.$$

Lemma 1 in [5] shows that the solution to the above optimization problem is given by  $\hat{Z} = \mathbf{S}_\lambda(W)$  where  $\mathbf{S}_\lambda(W) = U D_\lambda V'$  with  $D_\lambda = \text{diag}\{(d_1 - \lambda)_+, \dots, (d_r - \lambda)_+\}$ ,  $U D V'$  is the SVD of  $W$ ,  $D = \text{diag}(d_1, \dots, d_r)$ , and  $t_+ = \max(t, 0)$ .  $\mathbf{S}_\lambda(W)$  is a type of soft-thresholding operator.

As we have shown, Soft-Impute Algorithm has a convergence rate of  $O\left(\frac{1}{k}\right)$ , which for large  $k$  is worse than the Accelerated Soft-Impute Algorithm with rate  $O\left(\frac{1}{k^2}\right)$ . However, the numeric results in [5] show that Soft-Impute Algorithm performs favorably compared with the accelerated version illustrated in Figure 5 of [5].

## References

- [1] Amir Beck and Marc Teboulle. A fast iterative shrinkage-thresholding algorithm for linear inverse problems. *SIAM journal on imaging sciences*, 2(1):183–202, 2009.
- [2] Ronald E Bruck Jr. An iterative solution of a variational inequality for certain monotone operators in hilbert space. *Bulletin of the American Mathematical Society*, 81(5):890–892, 1975.
- [3] Nathan Halko, Per-Gunnar Martinsson, and Joel A Tropp. Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions. *SIAM review*, 53(2):217–288, 2011.
- [4] Trevor Hastie, Robert Tibshirani, and Martin Wainwright. *Statistical learning with sparsity*. CRC press, 2015.
- [5] Rahul Mazumder, Trevor Hastie, and Robert Tibshirani. Spectral regularization algorithms for learning large incomplete matrices. *Journal of machine learning research*, 11(Aug):2287–2322, 2010.
- [6] Yurii Nesterov et al. Gradient methods for minimizing composite objective function, 2007.
- [7] Tae-Hyun Oh, Yasuyuki Matsushita, Yu-Wing Tai, and In So Kweon. Fast randomized singular value thresholding for nuclear norm minimization. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 4484–4493, 2015.
- [8] Raymond KW Wong and Thomas Lee. Matrix completion with noisy entries and outliers. *arXiv preprint arXiv:1503.00214*, 2015.

# APPENDIX

## C code via LAPACK

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #define Epsilon 0.001
5 #define Iteration 10000
6
7 void dgesvd_(char *jobu, char *jobvt, int *m, int *n, double *a, int *lda, double *s, double
   *u,
8           int *ldu, double *vt, int *ldvt, double *work, int *lwork, int *info);
9 void dgemm_(char *transa, char *transb, int *m, int *n, int *k, double *alpha, double *a,
   int *lda,
10           double *b, int *ldb, double *beta, double *c, int *ldc);
11
12 int main(int argc, char *argv[]) {
13     //The 1st argument is the file name; the 2nd argument is lambda
14     char *filename = argv[1];
15     FILE *f = fopen(filename, "r");
16     double lambda = atof(argv[2]);
17     int N, P, i=0, j, id=0;
18
19     //In the dataset, the first 2 observations are N and P
20     fscanf(f, "%d", &N);
21     fscanf(f, "%d\n", &P);
22     double data[N*P], X[N*P], Zcomp[N*P], Ztemp[N*P], Znew[N*P], Zold[N*P];
23
24     // Read in the dataset
25     while (!feof(f)) {
26         fscanf(f, "%lf*~lf", &data[i]);
27         i++;
28     }
29
30     //Transpose(data) gives X
31     for (j=0; j<P; j++) {
32         for (i=0; i<N; i++) {
33             X[j*N+i] = data[i*P+j];
34         }
35     }
36     //Initialize Zold with 0
37     for (i=0; i>N*P; i++) {
38         Zold[i] = 0;
39     }
40
41     while (id < Iteration) {
42         //Compute the complementary matrix of Zold
43         for (i=0; i<N*P; i++) {
44             if (X[i]==0) {
45                 Zcomp[i]=Zold[i];
46             } else {
47                 Zcomp[i] = 0;
48             }
49         }
50
51         //Add X and Zcomp up
52         for (i=0; i<N*P; i++) {
53             Ztemp[i] = X[i] + Zcomp[i];
54         }
55
56         char jobu='A', jobvt='A';
57         double S[P], U[N*N], VT[P*P], work[5*N];
58         int lwork=5*N, info;
59         // Do SVD to Ztemp
60         dgesvd_(&jobu, &jobvt, &N, &P, Ztemp, &N, S, U, &N, VT, &P, work, &lwork, &info);
```



```

61 // Compute Dlambda
62 for (i=0; i<P; i++) {
63     if (S[i]-lambda>0) {
64         S[i] = S[i] - lambda;
65     } else {
66         S[i] = 0;
67     }
68 }
69 }
70 //Compute Znew
71 double UD[N*P], alpha=1.0, beta=0.0;
72 char transa='N', transb='N';
73 for (j=0; j<P; j++) {
74     for (i=0; i<N; i++) {
75         UD[N*j+i] = U[N*j+i] * S[j];
76     }
77 }
78 dgemm_(&transa, &transb, &N, &P, &P, &alpha, UD, &N, VT, &P, &beta, Znew, &N);
79
80
81 //Compare the ratio and Epsilon
82 double Nnorm=0, Dnorm=0, ratio;
83 for (i=0; i<N*P; i++) {
84     Nnorm = Nnorm + pow(Znew[i]-Zold[i],2);
85     Dnorm = Dnorm + pow(Zold[i], 2);
86 }
87 ratio = Nnorm/Dnorm;
88 if (ratio<Epsilon) break;
89 //Update Zold with Znew
90 for (i=0; i<N*P; i++) {
91     Zold[i] = Znew[i];
92 }
93 id++;
94 }
95 //Print out the final result
96 for (i=0; i<N; i++) {
97     for (j=0; j<P; j++) {
98         printf("%lf", Znew[j*N+i]);
99     }
100     printf("\n");
101 }
102
103 return 0;
104 }

```

../Codes/Min/SoftImpute.c

## R internal svd function

```

1 soft_thres <- function(x,y){
2     ifelse((x-y)>0, x-y, rep(0,length(x)))
3 }
4
5 softimpute_internalsvd <- function(mat_train, lambda, thres_ratio = 10^(-3), mat_validation
6     = NULL){
7     starting_time <- Sys.time()
8     na_index <- is.na(mat_train)
9     Z_old <- mat_train
10    Z_old[na_index] <- 0
11    ratio <- 1 + thres_ratio
12    while(ratio>thres_ratio){
13        mat_train[na_index] <- Z_old[na_index]
14        svd_result <- svd(mat_train)
15        Z_new <- svd_result$u%*%diag(soft_thres(svd_result$d, lambda))%*%t(svd_result$v)
16        ratio <- sum((Z_old-Z_new)^2)/sum((Z_old)^2)

```

```

16     #print(ratio)
17     Z_old <- Z_new
18 }
19 end_time <- Sys.time()
20 if(is.null(mat_validation))
21     return(list(Z_new, end_time - starting_time))
22 else{
23     return(c(sqrt(mean((mat_validation - Z_new)[!is.na(mat_validation)]^2)),
24             as.numeric(end_time - starting_time, units = "secs")))
25 }
26 }

```

../Codes/haozhe/softimpute\_internalsvd.R

## svd package

```

1 library(svd)
2
3 soft_thres <- function(x,y){
4     ifelse((x-y)>0, x-y, rep(0,length(x)))
5 }
6
7 softimpute_rpacksvd <- function(mat_train, lambda, thres_ratio = 10^(-3), mat_validation =
8     NULL){
9     starting_time <- Sys.time()
10    na_index <- is.na(mat_train)
11    Z_old <- mat_train
12    Z_old[na_index] <- 0
13    ratio <- 1 + thres_ratio
14    while(ratio>thres_ratio){
15        mat_train[na_index] <- Z_old[na_index]
16        svd_result <- propack.svd(mat_train)
17        Z_new <- svd_result$u%%diag(soft_thres(svd_result$d, lambda))%*%t(svd_result$v)
18        ratio <- sum((Z_old-Z_new)^2)/sum((Z_old)^2)
19        #print(ratio)
20        Z_old <- Z_new
21    }
22    end_time <- Sys.time()
23    if(is.null(mat_validation))
24        return(list(Z_new, end_time - starting_time))
25    else{
26        return(c(sqrt(mean((mat_validation - Z_new)[!is.na(mat_validation)]^2)),
27                as.numeric(end_time - starting_time, units = "secs")))
28    }
29 }

```

../Codes/haozhe/softimpute\_rpacksvd.R

## RcppArmadillo Package

## irlba package